DATA 298B

MSDA Project II Workbook 1

Team 3

Aparna S, Pavan S, Sanjana C, Shravani G, Sujata J

4. Model Development

4.1 Model Proposals

MedLlama3

MedLlama3 is a proposed adaptation of the Llama-3 language model, specifically engineered to address the unique demands of healthcare applications, as conceptualized by Touvron (2023). This model aims to serve as a powerful tool for processing advanced medical queries, delivering responses that are both coherent and medically accurate. By leveraging a deep understanding of healthcare-related topics, MedLlama3 is designed to support a range of digital health applications, including preliminary patient consultations, symptom analysis, and patient education. Its development is rooted in the need for reliable, context-aware language models capable of enhancing healthcare delivery through technology. As one of the proposed models for my project, MedLlama3 is positioned to contribute to accessible and efficient health solutions.

Architecture Overview

MedLlama3 is built upon the transformer-based architecture of Llama-3, as illustrated in Figure 64 (sourced from Google). This architecture is optimized for handling complex text-processing tasks, making it well-suited for medical applications that require precise and contextually relevant outputs. The model's design enables it to interpret and generate natural language by processing input text through multiple layers of transformation, ensuring that medical terminology and context are accurately captured.

Figure

Llama3 Architecture

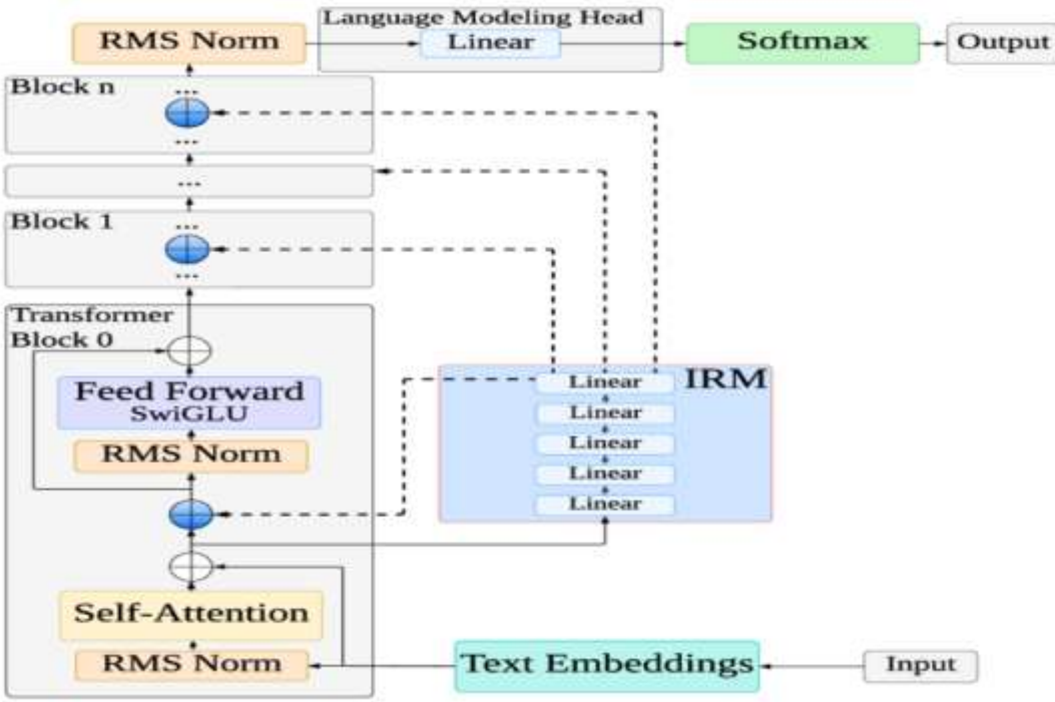

The architecture of MedLlama3 is composed of several key components that work together to process and generate medically relevant text:

1. **Text Embeddings**
   - Input text, such as medical queries or patient data, is converted into numerical embeddings.
   - These embeddings encode the semantic meaning of words, capturing the nuances of medical terminology and context.
   - This initial transformation ensures that the model can interpret complex healthcare-related language effectively.

2. **Transformer Blocks**
   - A stack of transformer blocks (Block 0 to Block n) processes the embeddings through iterative layers.
   - Each block consists of the following sub-components:

- **Self-Attention Mechanism**: Enables the model to focus on relationships between words within the input text, capturing long-range dependencies critical for understanding medical context (e.g., linking symptoms to potential conditions).
- **Feed Forward Layer (SwiGLU)**: Applies non-linear transformations to refine the representations, enhancing the model's ability to process intricate medical data.
- **RMS Normalization**: Stabilizes the training process by normalizing the input scale for each layer, ensuring consistent and reliable performance.

- The iterative processing across multiple blocks allows MedLlama3 to build a deep understanding of the input text.

3. **Intermediate Representation Module (IRM)**
   - Comprises a series of linear layers that act as an intermediary between transformer blocks.
   - The IRM preserves and transforms learned representations, ensuring that critical medical context is maintained throughout the processing pipeline.
   - This module enhances the model's ability to handle complex, multi-step reasoning tasks in healthcare scenarios.
4. **Language Modeling Head and Softmax**
   - The final layer maps the processed representations back to the model's vocabulary space.
   - The Softmax function generates probabilities for each token, enabling the model to produce coherent and medically accurate text outputs.
   - This component is crucial for generating responses that are both understandable and relevant to healthcare queries.

Functional Capabilities

The transformer-based architecture of MedLlama3 enables it to perform the following tasks effectively:

- **Query Processing**: Accurately interpret and respond to advanced medical questions, such as those related to diagnoses, treatments, or medical research.
- **Contextual Understanding**: Capture relationships between medical terms and concepts, ensuring responses are contextually appropriate.
- **Text Generation**: Produce clear, concise, and medically relevant text for applications like patient education materials or consultation summaries.
- **Scalability**: Handle a wide range of healthcare-related tasks, from symptom assessment to providing informational resources.

This architecture ensures that MedLlama3 can process complex medical texts with high accuracy, making it a robust foundation for healthcare application

MedLlama3 is proposed as a key model for my project due to its specialized focus on healthcare and its robust architectural foundation. The transformer-based design, with its ability to capture contextual relationships and process complex medical texts, aligns with the project's goal of developing advanced digital health solutions. MedLlama3's potential to support preliminary consultations, symptom evaluations, and patient education makes it a versatile tool for improving healthcare accessibility and efficiency. By integrating MedLlama3, the project can leverage cutting-edge language modeling to address real-world healthcare challenges, ensuring that responses are both accurate and actionable.

Potential Applications

- **Preliminary Consultations**: Assist healthcare providers by offering initial assessments based on patient queries, streamlining the consultation process.
- **Symptom Analysis**: Evaluate patient-reported symptoms to suggest possible conditions or recommend further medical evaluation.
- **Patient Education**: Generate accessible, medically accurate educational content to empower patients with knowledge about their health.
- **Research Support**: Aid medical professionals in accessing and summarizing relevant literature or data for clinical decision-making.

BioMistral 7B

BioMistral 7B is an open-source large language model (LLM) specifically designed for biomedical and healthcare applications, as introduced by Labrak et al. (2024). Built upon the Mistral 7B foundation model, BioMistral 7B has been further pre-trained on a vast corpus of biomedical literature from PubMed Central, enabling it to develop a deep understanding of medical terminology, concepts, and relationships. This model is tailored to address the unique challenges of adapting general-purpose LLMs to the medical domain, offering superior performance in medical question-answering (QA) tasks and other healthcare-related applications. As a proposed model for my project, BioMistral 7B aims to enhance digital health solutions by providing accurate, contextually relevant responses for tasks such as medical research, preliminary diagnostics, and patient education. Its open-source nature, released under the Apache 2.0 license, fosters collaboration and innovation in medical AI, making it a valuable asset for advancing healthcare technologies.

Architecture Overview

BioMistral 7B leverages the transformer-based architecture of Mistral 7B, enhanced with specialized pre-training and optimization techniques to excel in medical applications. The model incorporates advanced mechanisms such as grouped-query attention (GQA) and sliding window attention (SWA), which improve inference speed and enable efficient handling of long sequences, respectively. Its architecture is designed to process complex medical texts with high accuracy, making it suitable for tasks requiring precise and contextual understanding.

Figure

Mistral 7B Architecture

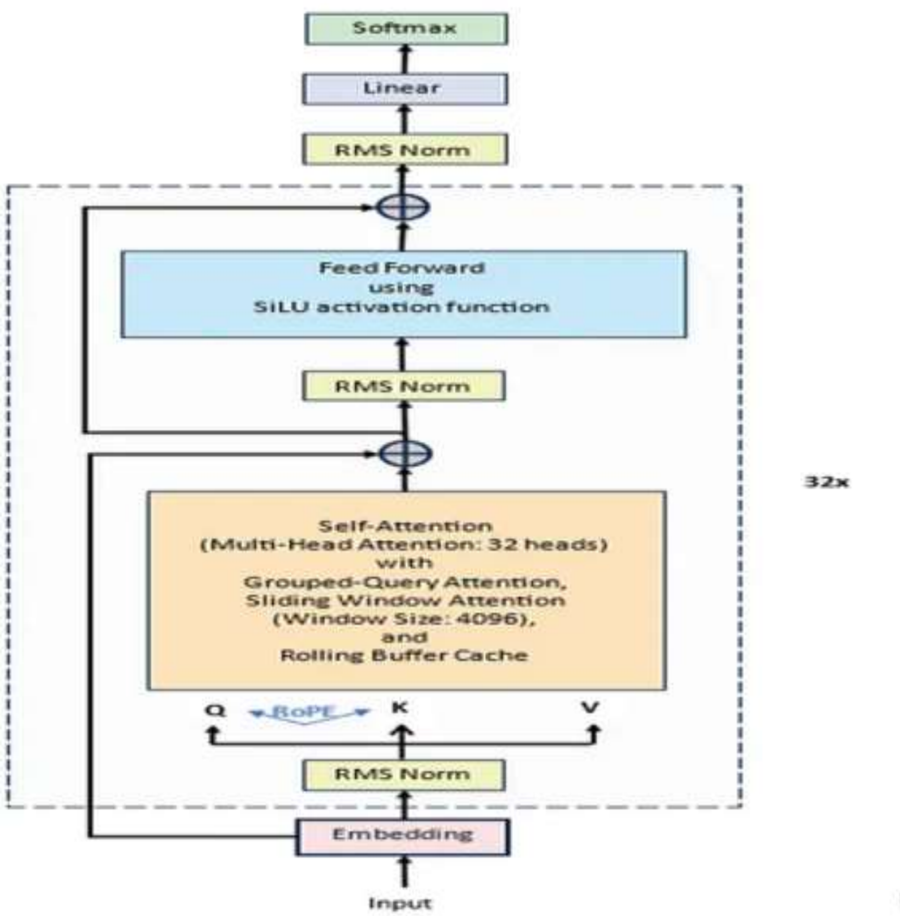

The architecture of BioMistral 7B is composed of several key components that enable it to process and generate medically relevant text effectively:

1. **Text Embeddings**

- Input text, such as medical queries or clinical notes, is transformed into numerical embeddings that capture the semantic meaning of words and phrases.
- These embeddings are optimized for biomedical terminology, ensuring accurate representation of complex medical concepts.
- The embedding layer serves as the foundation for contextual understanding in healthcare applications.

2. **Transformer Blocks**

- A stack of transformer blocks (Block 0 to Block n) processes the embeddings through multiple layers.
- Each block includes:
    - **Self-Attention Mechanism (with Grouped-Query Attention)**: Captures relationships between words across the input text, enabling the model to understand long-range dependencies critical for medical context (e.g., linking symptoms to diagnoses). GQA enhances inference speed by grouping queries, reducing computational overhead.
    - **Feed Forward Layer (SwiGLU)**: Applies non-linear transformations to refine representations, improving the model's ability to process intricate medical data.
    - **RMS Normalization**: Stabilizes training by normalizing the input scale for each layer, ensuring consistent performance across diverse medical tasks.
- The iterative processing across transformer blocks builds a deep, contextual understanding of the input.

3. **Sliding Window Attention (SWA)**

- Enables efficient handling of long sequences by limiting attention to a fixed window of previous tokens (e.g., 4,096 hidden states).
- Reduces memory usage and inference costs, making the model suitable for processing lengthy medical documents or patient records.
- SWA ensures that BioMistral 7B can maintain performance on extended inputs without compromising quality.

4. **Intermediate Representation Module (IRM)**

- Comprises linear layers that preserve and transform learned representations between transformer blocks.
- Ensures that critical medical context is maintained throughout the processing pipeline, supporting complex reasoning tasks.
- Enhances the model’s adaptability to specialized biomedical applications.

5. **Language Modeling Head and Softmax**

- The final layer maps processed representations back to the vocabulary space, generating probabilities for each token.
- The Softmax function produces coherent and medically accurate text outputs, suitable for answering queries or generating reports.
- This component ensures that responses are both understandable and relevant to healthcare contexts.

Optimization Techniques

BioMistral 7B incorporates several optimization strategies to enhance performance and efficiency:

- **Pre-Training on PubMed Central**: Extensive training on biomedical literature equips the model with domain-specific knowledge, improving its accuracy in medical tasks.

- **Quantization**: Lightweight variants (e.g., 4-bit, 8-bit) reduce memory footprint (VRAM requirements from 4.68GB to 15.02GB) and improve inference speed, making the model deployable on consumer-grade devices.
- **Model Merging**: Techniques such as SLERP, TIES, and DARE combine BioMistral's biomedical expertise with Mistral's general knowledge, creating hybrid models that balance specialized and broad capabilities.
- **Supervised Fine-Tuning (SFT)**: Fine-tuning on medical QA datasets enhances accuracy, with BioMistral achieving an average accuracy of 57.3% across 10 medical QA benchmarks, outperforming other open-source medical models.

Functional Capabilities

The architecture and optimizations enable BioMistral 7B to perform the following tasks effectively:

- **Medical Question-Answering**: Provide accurate responses to queries about conditions, treatments, and medical research, as demonstrated by strong performance on benchmarks like Clinical KG and MedQA.
- **Text Generation**: Create coherent medical text, such as educational content or clinical summaries, for healthcare professionals and patients.
- **Text Classification**: Identify medical conditions or categorize clinical texts based on input data.
- **Multilingual Processing**: Evaluated on seven languages, supporting global medical research and applications.

- **Research Support**: Facilitate analysis of medical literature, aiding researchers in summarizing and extracting insights.

This architecture, combined with domain-specific training and optimizations, positions BioMistral 7B as a leading open-source model for biomedical applications.

BioMistral 7B is proposed as a key model for my project due to its specialized design for healthcare and its competitive performance against both open-source and proprietary medical LLMs. Its transformer-based architecture, enhanced by GQA and SWA, ensures efficient and accurate processing of complex medical texts, aligning with the project's goal of developing advanced digital health solutions. The model's open-source availability under the Apache 2.0 license encourages collaborative development, while its quantization and model merging capabilities make it adaptable to resource-constrained environments. BioMistral 7B's strong performance on medical QA benchmarks (average accuracy of 57.3%, with variants reaching 59.4%) and its multilingual capabilities further enhance its suitability for global healthcare applications. By integrating BioMistral 7B, the project can leverage cutting-edge AI to improve healthcare accessibility, support medical research, and deliver reliable patient-facing tools.

Potential Applications

- **Medical Question-Answering**: Answer queries about symptoms, treatments, or medical conditions, supporting healthcare professionals and patients.
- **Preliminary Diagnostics**: Assist in symptom analysis to suggest potential conditions, streamlining initial assessments.
- **Patient Education**: Generate accessible, medically accurate content to inform patients about their health.

- **Medical Research**: Summarize and analyze biomedical literature, aiding researchers in drug discovery and clinical studies.
- **Clinical Documentation**: Support healthcare providers by generating structured summaries or reports from patient data.
- **Multilingual Support**: Facilitate medical research and patient care in diverse linguistic contexts, leveraging its evaluation across seven languages.

BioMistral 7B represents a significant advancement in open-source medical AI, combining the robust Mistral 7B architecture with specialized biomedical training to address healthcare challenges. Its efficient design, strong benchmark performance, and open-source accessibility make it an ideal candidate for my project's mission to enhance digital health solutions. By leveraging BioMistral 7B, the project can deliver accurate, context-aware tools for medical research, diagnostics, and patient education, contributing to the broader adoption of AI-driven healthcare innovations

Meditron

Meditron is a suite of open-source large language models (LLMs) specifically tailored for the medical domain, developed by researchers at École Polytechnique Fédérale de Lausanne (EPFL), Yale School of Medicine, and supported by the International Committee of the Red Cross (ICRC) (Chen et al., 2023). Built upon the Llama-2 and Llama-3.1 architectures, Meditron includes models with 7B and 70B parameters (Meditron-7B and Meditron-70B) and a newer variant, Llama-3.1-Meditron-3 (8B and 70B), designed to democratize access to medical knowledge. By leveraging domain-adaptive pre-training on a curated medical corpus, Meditron excels in medical reasoning tasks, outperforming other open-source models and even some closed-source commercial LLMs like GPT-3.5 and Med-PaLM on standard benchmarks. Its applications include supporting clinical decision-making, medical question-answering, differential diagnosis, and providing evidence-based health information, particularly in low-resource and humanitarian settings. As a proposed model for my project, Meditron's open-source nature, high performance, and focus on equitable healthcare access make it a compelling candidate for advancing digital health solutions.

Architecture Overview

Meditron is a transformer-based model adapted from Llama-2 (for Meditron-7B and 70B) and Llama-3.1 (for Meditron-3), optimized for medical applications through extensive pre-training on a 48.1B-token medical corpus called GAP-Replay. This corpus includes clinical guidelines, PubMed abstracts, full-text medical papers, and general-domain data from RedPajama-v1, ensuring a robust and diverse knowledge base. The architecture is designed for efficiency and scalability, utilizing distributed training techniques and a sophisticated structure to handle complex medical texts.

Meditron's architecture is a causal decoder-only transformer, based on Llama-2 (7B: 32 layers, 32 attention heads, 4096 hidden dimensions; 70B: 64 layers, 64 attention heads, 8192 hidden dimensions) and Llama-3.1, optimized for medical text processing:

1. **Text Embeddings**
    - Input text (e.g., medical queries, clinical notes) is converted into numerical embeddings that encode semantic meaning.
    - Embeddings are tailored to capture medical terminology and context, leveraging pre-training on biomedical data.
    - Supports text-only input with a maximum sequence length of 2,048 tokens.
2. **Transformer Blocks**

- A stack of transformer blocks (32 for 7B, 64 for 70B) processes embeddings iteratively. Each block includes:
    - **Self-Attention Mechanism**: Captures long-range dependencies in text, critical for understanding relationships in medical contexts (e.g., symptoms to diagnoses). Uses multi-head attention (32 heads for 7B, 64 for 70B) for robust contextualization.
    - **Feed Forward Layer (SwiGLU)**: Applies non-linear transformations to refine representations, enhancing the model's ability to process complex medical data.
    - **RMS Normalization**: Stabilizes training by normalizing layer inputs, improving performance across diverse tasks.
- Blocks are optimized for distributed training using Megatron-LLM, enabling efficient scaling.

3. **Intermediate Representation Module (IRM)**
    - Comprises linear layers that preserve and transform representations between transformer blocks.
    - Maintains critical medical context, supporting multi-step reasoning in clinical scenarios.
    - Enhances adaptability to specialized medical tasks.
4. **Language Modeling Head and Softmax**
    - Maps processed representations back to the vocabulary space, generating token probabilities.
    - The Softmax function produces coherent, medically accurate text outputs for tasks like question-answering or report generation.
    - Optimized for high-quality text generation in medical contexts.

Optimization Techniques

Meditron incorporates several optimizations to enhance performance and accessibility:

- **Domain-Adaptive Pre-Training (GAP-Replay)**: Trained on 48.1B tokens, including 46K clinical guidelines, 16.1M PubMed abstracts, 5M full-text medical papers, and 400M general-domain tokens, ensuring comprehensive medical knowledge.
- **Distributed Training**: Utilizes Megatron-LLM on 16 nodes with 8 NVIDIA A100 (80GB) GPUs, achieving a throughput of ~40,200 tokens/second.
- **In-Context Learning**: Supports downstream tasks with 3–5 demonstrations in prompts, improving adaptability without fine-tuning.
- **Fine-Tuning Flexibility**: Can be fine-tuned, instruction-tuned, or RLHF-tuned for specific tasks like medical QA or diagnosis support.
- **Meditron-V (Multimodal Variant)**: Extends capabilities to process biomedical imaging (e.g., histology, radiology), outperforming models like Med-PaLM M (562B) on multimodal benchmarks.

Functional Capabilities

Meditron's architecture and training enable it to perform the following tasks:

- **Medical Question-Answering**: Achieves high accuracy on benchmarks like MedQA (77.6% for 70B), PubMedQA, and MedMCQA, outperforming Llama-2-70B and GPT-3.5.
- **Clinical Decision Support**: Assists with differential diagnosis and disease information queries (e.g., symptoms, causes, treatments).

- **Text Generation**: Produces coherent medical text for educational content or clinical summaries.
- **Multimodal Reasoning (Meditron-V)**: Processes medical imaging alongside text for tasks like radiology report generation.
- **Humanitarian Applications**: Incorporates ICRC clinical guidelines, supporting low-resource and diverse healthcare contexts.

This architecture, combined with domain-specific pre-training, positions Meditron as a leading open-source medical LLM suite.

Meditron is proposed for my project due to its exceptional performance, open-source accessibility, and alignment with the goal of advancing equitable healthcare solutions. Its ability to outperform closed-source models like GPT-3.5 and Med-PaLM on medical benchmarks (e.g., 6% absolute gain over Llama-2 baselines, within 5% of GPT-4 on MedQA) demonstrates its robustness for medical tasks. The open-source release of model weights, training code, and the GAP-Replay corpus fosters transparency and enables further customization, critical for research and development. Meditron's focus on humanitarian contexts, supported by ICRC guidelines, ensures relevance in low-resource settings, addressing underrepresented populations and neglected diseases. The Meditron MOOVE initiative, which invites global healthcare professionals to validate performance in real-world scenarios, further enhances its potential for practical impact. By integrating Meditron, my project can leverage a scalable, transparent, and medically specialized LLM to improve clinical decision-making, patient education, and medical research.

Potential Applications

- **Medical Question-Answering:** Provide accurate responses to queries about conditions, treatments, or medical guidelines, supporting clinicians and patients.

- **Differential Diagnosis**: Assist healthcare providers in identifying potential conditions based on symptoms, enhancing diagnostic workflows.
- **Patient Education**: Generate accessible, evidence-based content to inform patients about health conditions and treatments.
- **Clinical Documentation**: Support generation of structured summaries or reports from patient data, streamlining workflows.
- **Medical Research**: Summarize and analyze PubMed literature, aiding researchers in drug discovery or clinical studies.
- **Humanitarian Healthcare**: Support clinical decision-making in low-resource settings, leveraging ICRC guidelines for equitable care.
- **Multimodal Applications (Meditron-V)**: Analyze medical imaging alongside text for tasks like radiology or histology interpretation.

Meditron represents a groundbreaking advancement in open-source medical AI, combining a scalable transformer-based architecture with domain-specific pre-training to deliver state-of-the-art performance in medical reasoning tasks. Its open-source availability, high accuracy (e.g., 77.6% on MedQA for 70B), and focus on humanitarian and low-resource settings make it an ideal candidate for my project's mission to enhance digital health solutions. By leveraging Meditron, the project can develop transparent, equitable, and evidence-based tools for clinical decision-making, patient education, and medical research, contributing to the global democratization of healthcare knowledge.

BioGPT

BioGPT, developed by Microsoft Research (Luo et al., 2022), is a generative large language model (LLM) specifically designed for biomedical applications. Built upon the GPT architecture, BioGPT is pre-trained on a large corpus of 15 million PubMed abstracts, encompassing approximately 2 billion tokens, to capture the nuances of biomedical language and concepts. With model sizes ranging from 347M (BioGPT) to 1.5B (BioGPT-Large) parameters, and a 2.7B instruction-tuned variant (BioGPT-JSL), it excels in generating coherent and contextually relevant medical text, making it ideal for tasks such as medical question-answering, patient education, and clinical text generation. BioGPT's autoregressive design enables it to produce fluent, natural language outputs, positioning it as a powerful tool for digital health applications. As a proposed model for my project, BioGPT's open-source availability, generative capabilities, and strong performance in biomedical tasks make it a valuable addition to enhance healthcare accessibility and support medical research.

Architecture Overview

BioGPT is a transformer-based, autoregressive model optimized for text generation in the biomedical domain. Its architecture, derived from GPT, leverages a decoder-only structure with enhancements tailored for medical text processing. The model's design prioritizes fluency and scalability, enabling it to handle a wide range of medical tasks efficiently.

BioGPT's architecture is composed of key components that enable it to process and generate medically relevant text:

1. **Text Embeddings**

    - Input text, such as medical queries or research abstracts, is transformed into numerical embeddings that encode the semantic meaning of words and phrases.
    - Embeddings are optimized for biomedical terminology, ensuring accurate representation of complex medical concepts like diseases, drugs, and procedures.
    - This layer forms the foundation for generating contextually appropriate outputs.

2. **Transformer Blocks**

    - A stack of transformer blocks (e.g., 24 layers for BioGPT, 36 for BioGPT-Large) processes embeddings through iterative layers. Each block includes:

        - **Self-Attention Mechanism (Causal)**: Captures relationships between preceding tokens in the sequence, enabling the model to generate coherent text by focusing on prior context. Multi-head attention (e.g., 20 heads for BioGPT) enhances contextual understanding.
        - **Feed Forward Layer**: Applies non-linear transformations to refine representations, improving the model's ability to handle intricate medical data.
        - **Layer Normalization**: Stabilizes training by normalizing layer inputs, ensuring consistent performance across diverse tasks.

    - The autoregressive nature of the blocks ensures that text is generated sequentially, ideal for tasks requiring fluency.

3. **Intermediate Representation Module (IRM)**

- Comprises linear layers that maintain and transform learned representations between transformer blocks.
- Preserves critical medical context during generation, supporting tasks that require sustained coherence, such as summarizing research papers.
- Enhances the model's ability to produce structured outputs for specific tasks.

4. **Language Modeling Head and Softmax**

   - Maps processed representations back to the vocabulary space, generating probabilities for each token.
   - The Softmax function produces fluent and medically accurate text outputs, suitable for answering queries or generating educational content.
   - Optimized for high-quality text generation in biomedical contexts.

Optimization Techniques

BioGPT incorporates several optimizations to enhance performance and usability:

- **Pre-Training on PubMed Abstracts**: Trained on 15M PubMed abstracts (~2B tokens), providing deep domain-specific knowledge for biomedical tasks.
- **Instruction Tuning (BioGPT-JSL)**: Fine-tuning with instruction-based datasets improves adaptability for tasks like question-answering and text generation, achieving state-of-the-art results.
- **Quantization**: Lightweight variants (e.g., 4-bit or 8-bit) reduce memory requirements (e.g., ~700MB for BioGPT with quantization), enabling deployment on consumer-grade hardware.

- **Task-Specific Fine-Tuning**: Supports fine-tuning for tasks like relation extraction or QA, with strong results on benchmarks like PubMedQA.

Functional Capabilities

BioGPT's architecture and training enable it to perform the following tasks effectively:

- **Medical Question-Answering**: Achieves high accuracy on PubMedQA (78.2% for BioGPT-Large), providing detailed and fluent responses to medical queries.
- **Text Generation**: Produces coherent medical text for applications like patient education materials, clinical summaries, or research abstracts.
- **Relation Extraction**: Sets benchmarks on datasets like BC5CDR (F1: 90.22%) and KD-DTI (F1: 44.76%) for end-to-end relation extraction.
- **Text Classification**: Supports tasks like identifying medical entities or classifying clinical texts, though less specialized than discriminative models.
- **Research Support**: Summarizes and generates insights from biomedical literature, aiding researchers in drug discovery or clinical studies.

This architecture, combined with domain-specific pre-training and instruction tuning, makes BioGPT a leading generative model for biomedical applications.

BioGPT is proposed for my project due to its strong generative capabilities, open-source accessibility, and alignment with the goal of advancing digital health solutions. Its ability to produce fluent, medically accurate text makes it particularly suitable for patient-facing applications, such as generating educational content or conversational responses for preliminary consultations. BioGPT's performance on medical QA benchmarks (e.g., 78.2% accuracy on

PubMedQA for BioGPT-Large) and its instruction-tuned variants (e.g., BioGPT-JSL) demonstrate its effectiveness for open-ended medical queries. The model's open-source availability on Hugging Face, coupled with quantization options, ensures it can be deployed in resource-constrained environments, complementing the more computationally intensive models in my project (e.g., Meditron-70B). By integrating BioGPT, the project can leverage a lightweight, generative LLM to enhance healthcare accessibility, support patient education, and facilitate medical research.

Potential Applications

- **Medical Question-Answering**: Provide detailed, conversational responses to queries about conditions, treatments, or medical research, supporting clinicians and - **Preliminary Diagnostics**: Generate hypotheses or summaries based on symptom descriptions, streamlining initial assessments.
- **Patient Education**: Create accessible, medically accurate content to inform patients about health conditions and treatments.
- **Clinical Documentation**: Generate structured summaries or reports from patient data, improving workflow efficiency.
- **Medical Research**: Summarize and generate insights from PubMed literature, aiding researchers in drug discovery or clinical studies.
- **Conversational Interfaces**: Power chatbots or virtual assistants for digital health platforms, enhancing user engagement.

BioGPT represents a powerful and accessible solution for generative biomedical tasks, combining a scalable GPT-based architecture with domain-specific pre-training to deliver fluent and accurate medical text. Its open-source availability, strong performance on benchmarks like PubMedQA, and lightweight deployment options make it an ideal candidate for my project's mission to enhance digital health solutions. By integrating BioGPT, the project can develop user-friendly, evidence-based tools for patient education, medical question-answering, and research support, contributing to the broader adoption of AI-driven healthcare innovations. Compared to other proposed models (MedLlama3, BioMistral 7B, Meditron), BioGPT offers a specialized focus on generative tasks, complementing their broader capabilities and ensuring a versatile model suite for the project.

4.2 Model Supports

The healthcare AI assistant for chronic disease management is powered by a sophisticated integration of Google Cloud Platform (GCP) services, specialized open-source large language models (LLMs), and a Retrieval-Augmented Generation (RAG) framework. This system leverages fine-tuned LLMs—**MedLlama3, BioMistral-7B, BioGPT, and Meditron**—to deliver accurate, contextually relevant, and medically reliable responses. The RAG framework, enhanced by the **Joint Medical LLM and Retrieval Training (JMLR)** methodology, integrates real-time medical knowledge from PubMed APIs and vectorized medical data stored in a **Pinecone vector database**. This architecture ensures efficient retrieval of up-to-date information, enabling the AI assistant to support patient engagement, automate routine queries, and enhance healthcare efficiency. Deployed on GCP with a user-friendly front-end interface, the system is designed for scalability, security, and continuous improvement, with performance monitoring and feedback loops to optimize model accuracy and patient outcomes.

Key Technologies and Tools

1. **Google Cloud Storage (GCS)**
    - **Purpose**: Serves as the primary storage for raw and processed healthcare datasets, including MIMIC-III/IV, iCliniq, PubMedQA, and MedDialog.
    - **Role in RAG**: Stores raw documents (e.g., clinical notes, research papers) and pre-processed embeddings required for RAG workflows, ensuring seamless integration with other GCP services.
    - **Benefits**: Provides scalable, secure, and centralized data lakes, facilitating efficient data management and retrieval for model training and inference.
2. **Pinecone Vector Database**
    - **Purpose**: A managed, high-performance vector database that stores and queries embedding vectors generated from medical texts, enabling fast similarity searches for RAG.
    - **Role in RAG**: Hosts vectorized representations of medical knowledge from PubMed, clinical guidelines, and internal datasets, optimized for semantic retrieval using the JMLR methodology.
    - **Benefits**: Offers low-latency, scalable vector search capabilities, ensuring the AI assistant retrieves relevant and up-to-date medical information to enhance response accuracy.
3. **Google Dataflow**
    - **Purpose**: Executes Extract-Load-Transform (ELT) pipelines to preprocess raw healthcare data, performing cleansing, normalization, feature extraction, and aggregation.

- **Role**: Transforms datasets like MIMIC-III/IV and iCliniq into formats suitable for model training and embedding generation, writing processed data to BigQuery or Pinecone.
- **Benefits**: Automates scalable data processing, ensuring high-quality inputs for machine learning workflows.

4. **BigQuery**
   - **Purpose**: Acts as a data warehouse for storing pre-processed, normalized datasets used for analytics, feature engineering, and model training.
   - **Role**: Enables large-scale querying and feature selection, supporting ad hoc analyses and scheduled training pipelines for LLMs.
   - **Benefits**: Provides fast, SQL-based access to structured data, enhancing the efficiency of model development and evaluation.
5. **Vertex AI**
   - **Purpose**: A unified platform for developing, fine-tuning, and deploying the LLMs (MedLlama3, BioMistral-7B, BioGPT, Meditron).
   - **Fine-Tuning**: Models are fine-tuned on healthcare-specific datasets (MIMIC-III/IV, iCliniq, PubMedQA, MedDialog) using efficient techniques like **Low-Rank Adaptation (LoRA)** to optimize performance for medical tasks.
   - **RAG Integration**: Implements the JMLR methodology to combine LLM inference with real-time retrieval from PubMed APIs and Pinecone, ensuring factual and contextually enriched responses.
   - **Embedding Generation**: Uses PubMed-BERT to generate high-quality embeddings for medical texts, stored in Pinecone for semantic search.

- **Benefits**: Streamlines model training, deployment, and RAG workflows, ensuring scalability and domain-specific accuracy.

6. **Cloud Run & Front-End Interface**
   - **Ascendancy**: Deploy the fine-tuned LLMs and RAG system in a serverless environment using Cloud Run, ensuring scalable, low-latency responses.
   - **Front-End**: A user-friendly interface built with HTML, CSS, and JavaScript allows patients and healthcare providers to interact seamlessly with the AI assistant.
   - **Benefits**: Provides an intuitive platform for chronic disease management, accessible via telemedicine platforms, hospital systems, or self-care applications.
7. **Cloud Logging and Cloud Monitoring**
   - **Purpose**: Tracks model performance, resource utilization, and response accuracy, capturing metrics like latency, error rates, and factual consistency.
   - **Role**: Enables continuous feedback loops, where real-world interactions are logged and analyzed to improve the training pipeline and model performance.
   - **Benefits**: Ensures system stability and supports data-driven optimization for enhanced patient outcomes.
8. **Security and Compliance**
   - **Encryption**: Implements AES-256 encryption to secure patient and medical data at rest and in transit.
   - **Access Control**: Uses role-based access control (RBAC) to enforce strict access policies, ensuring compliance with HIPAA regulations.
   - **Benefits**: Protects sensitive healthcare data, maintaining patient trust and regulatory adherence.

Overview of End-to-End Architecture

Figure: End-to-End Architecture

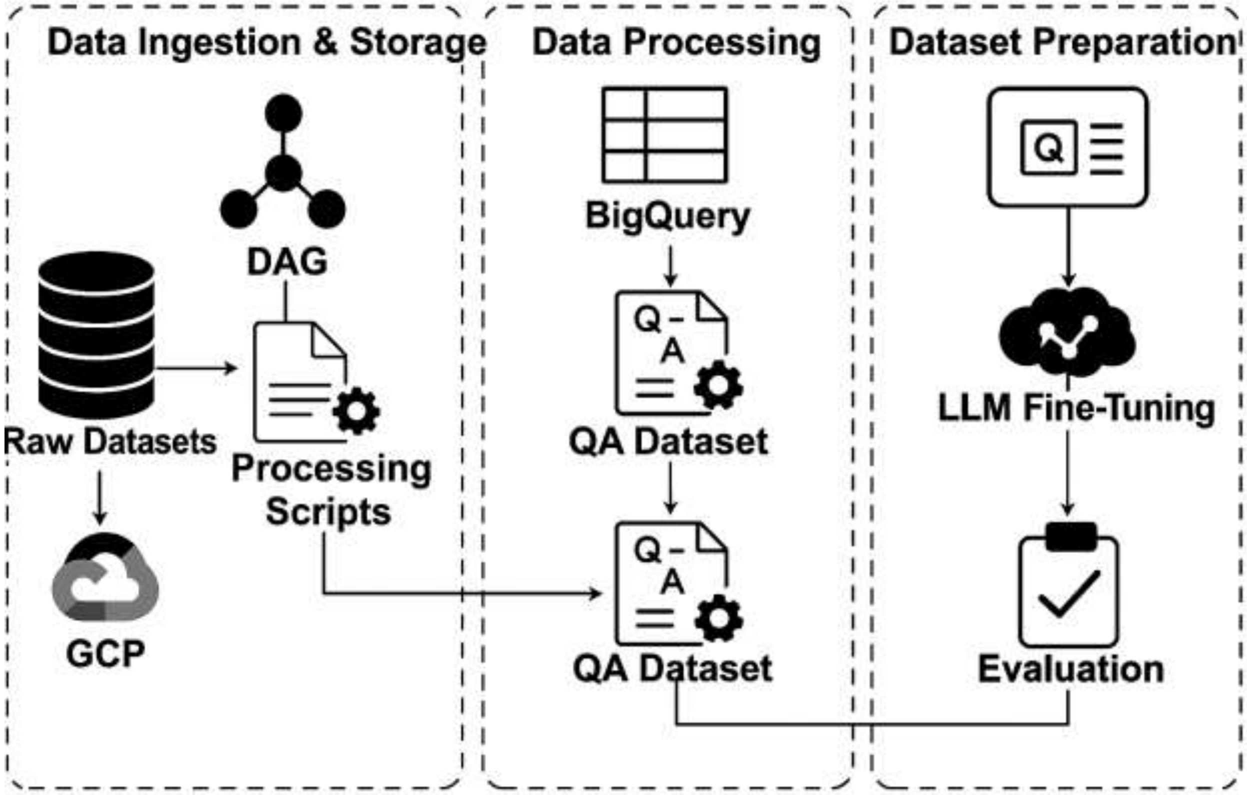

The figure illustrates the end-to-end data pipeline of the project. Raw medical datasets stored in GCP are processed using DAG-based scripts, and the cleaned outputs are stored in BigQuery. From there, a question-answer (QA) dataset is constructed and used to fine-tune large language models (LLMs), followed by evaluation to validate model performance.

Data Ingestion & Storage

- **Source Integration**: Supports ingestion from local folders, mounted Google Drive locations, and APIs (e.g., PubMed), enabling import of clinical datasets, logs, and structured forms.
- **Scheduled Batch Jobs**: Cloud Scheduler triggers batch jobs to extract and upload data to GCS at regular intervals.

- **GCS**: Centralizes raw and semi-structured datasets in scalable, secure data lakes.
- **Dataflow for Preprocessing**: Performs ELT operations—cleansing, normalization, and feature extraction—preparing data for analytics and ML workflows.
- **BigQuery for Analytics**: Loads processed data into BigQuery for large-scale querying, feature selection, and training pipelines.

Model Training and Knowledge Retrieval

- **Feature Engineering**: Preprocesses data for enhanced prediction precision and contextual understanding.
- **Embedding Generation**: Uses PubMed-BERT via Vertex AI to generate embeddings for semantic understanding, stored in Pinecone.
- **Domain-Specific Fine-Tuning**: Fine-tunes MedLlama3, BioMistral-7B, BioGPT, and Meditron on healthcare datasets using LoRA, specializing models for medical queries.
- **RAG with JMLR**: Dynamically retrieves context from PubMed APIs and Pinecone using the JMLR methodology, ensuring factual, up-to-date responses.

Model Deployment and Monitoring

- **Deployment**: Deploys fine-tuned LLMs via Vertex AI Model Garden, ensuring scalability and security.
- **Monitoring & Logging**: Cloud Monitoring tracks performance metrics (e.g., latency, usage patterns), while Cloud Logging captures user interactions for optimization.
- **Cloud Functions**: Structures user queries and routes them to the LLM for context-aware responses.

- **Front-End Interface**: A HTML/CSS/JavaScript-based UI facilitates seamless patient and provider interactions.
- **Feedback Loop**: Captures real-time feedback to optimize performance and relevance.

Model Retraining, Re-Deployment, and Monitoring

- **New Data Ingestion**: Periodically collects updated clinical and conversational data from GCS and BigQuery.
- **Fine-Tuning & Registry**: Fine-tunes LLMs with new data, registering models in Vertex AI Model Registry for traceability.
- **Evaluation & Versioning**: Evaluates models using BLEU, ROUGE, Medical Concept Recall (MCR), and Factual Consistency Score, maintaining version metadata for auditability.

Benefits and Alignment with Project Goals

This architecture supports the project's goal of improving chronic disease management by:

- **Enhancing Patient Engagement**: Provides on-demand, accurate medical information via an intuitive interface.
- **Reducing Healthcare Burden**: Automates' routine queries, minimizing unnecessary hospital visits.
- **Ensuring Accuracy**: Combines fine-tuned LLMs with RAG and JMLR for factual, context-aware responses.
- **Maintaining Compliance**: Implements AES-256 encryption and RBAC for HIPAA-compliant data security.

- **Supporting Scalability**: Leverages GCP's scalable infrastructure and Pinecone's efficient vector search for real-time performance.

By integrating MedLlama3, BioMistral-7B, BioGPT, and Meditron with a robust RAG framework, Pinecone vector database, and GCP services, the healthcare AI assistant delivers a patient-focused solution that enhances outcomes, improves efficiency, and democratizes access to medical knowledge.

Figure 24

Platform architecture, components, machine learning data flows.

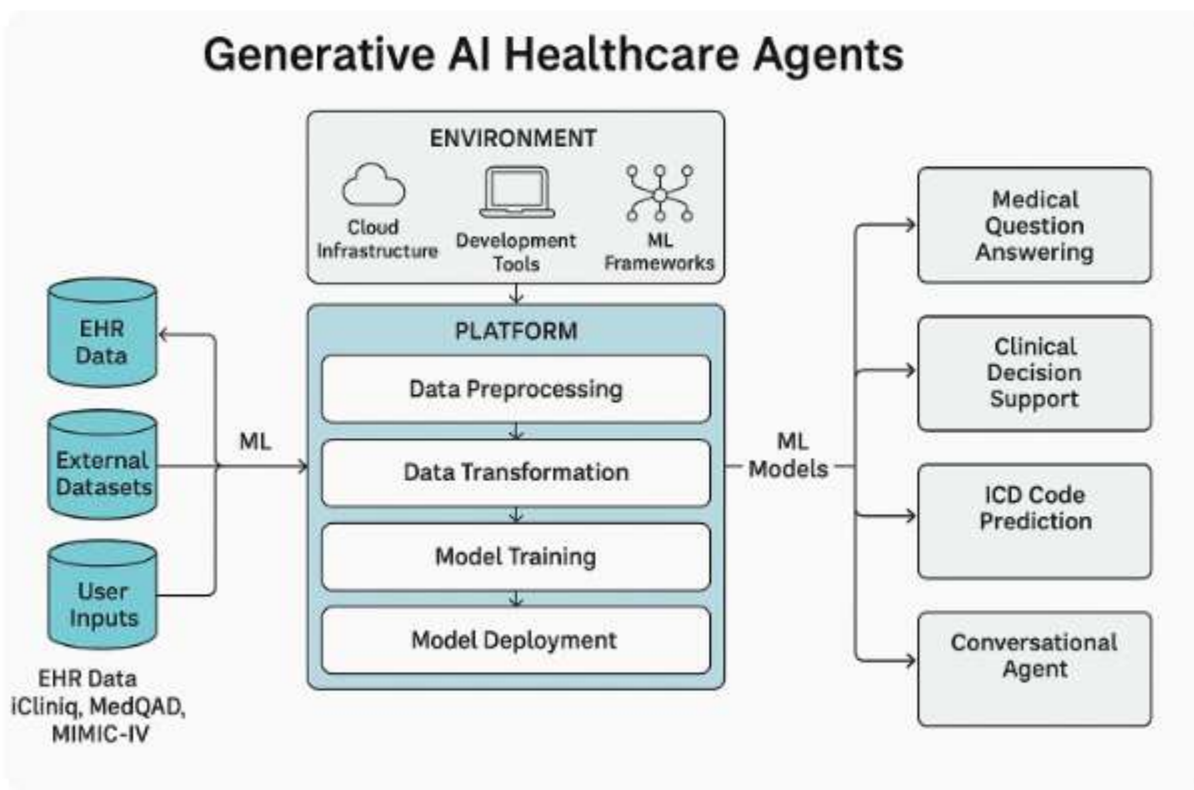

4.3Model Comparison and Justification

The healthcare AI assistant for chronic disease management integrates four open-source, healthcare-specialized large language models (LLMs): MedLlama3, BioMistral-7B, BioGPT, and

Meditron. These models were selected for their complementary strengths in addressing diverse requirements, including general medical question-answering, biomedical knowledge retrieval, generative text production, and high-performance medical reasoning. Each model is fine-tuned on datasets such as MIMIC-III/IV, iCliniq, PubMedQA, and MedDialog, and integrated with a Retrieval-Augmented Generation (RAG) framework using the Joint Medical LLM and Retrieval Training (JMLR) methodology. Vectorized medical knowledge is stored in a Pinecone vector database for efficient retrieval, ensuring accurate and contextually relevant responses. The following comparison table and justification outline the targeted problems, features, approaches, strengths, limitations, and data characteristics of each model, demonstrating their collective role in achieving the project's goals of enhancing patient engagement, reducing healthcare burden, and improving outcomes.

Model	Targeted Problem	Features	Approach	Strengths	Limitations	Data Size & Complexity	Types of Data
MedLlama3	General-purpose healthcare Q&A, patient education, preliminary health advice	Built on Llama-2; versatile in medical Q&A; handles complex medical terminology via self-attention mechanisms	Transformer-based deep learning with fine-tuning on healthcare datasets	High accuracy and versatility for diverse medical queries; ideal for general healthcare applications and patient-facing interactions	Computationally intensive; may require significant resources for real-time deployment in low-resource settings	High data size, complex queries	General medical queries, patient education materials, clinical notes
BioMistral-7B	Biomedical knowledge retrieval, clinical	Pre-trained on PubMed; domain-	NLP and deep learning with domain-	Exceptional in clinical knowledge	Limited real-time support due to computatio	Moderate to large, complex	Biomedical research data (e.g., PubMed), clinical

	research Q&A, specialized medical queries	specific language; supports multilingual queries; uses model-merging (SLERP, DARE)	adaptive pre-training	retrieval and biomedical terminology; excels in research-oriented and technical queries	nal complexity; less effective for general health queries	biomedical terms	guidelines, research papers
BioGPT	Generative medical text, patient education, conversational Q&A	Built on GPT; pre-trained on 15M PubMed abstracts; instruction-tuned variants (e.g., BioGPT-JSL)	Autoregressive transformer with instruction tuning and fine-tuning	Produces fluent, coherent medical text; ideal for patient education and open-ended QA; lightweight deployment with quantization	Unidirectional processing limits contextual depth for tasks like entity extraction; less effective for discriminative tasks	Moderate data size, moderate to complex queries	PubMed abstracts, patient education content, conversational queries
Meditron	Advanced medical reasoning, differential diagnosis, clinical decision support	Built on Llama-2/3.1; pre-trained on GAP-Replay (48.1B tokens); supports multimodal tasks (Meditron-V)	Transformer-based with domain-adaptive pre-training and in-context learning	High accuracy in medical QA (e.g., 77.6% on MedQA); supports complex reasoning and humanitarian applications; scalable (7B to 70B)	Research-only; requires extensive testing for clinical use; resource-intensive for 70B variant	Very high data size, complex scientific and diagnostic queries	Clinical guidelines, PubMed abstracts, full-text papers, structured medical data

Justification for Model Selection

The selection of **MedLlama3, BioMistral-7B, BioGPT, and Meditron** is driven by their complementary capabilities, which collectively address the diverse needs of a patient-focused healthcare AI assistant for chronic disease management. Each model contributes unique strengths to the system, ensuring robust performance across general healthcare queries, specialized biomedical tasks, generative text production, and advanced medical reasoning. Their integration with a RAG framework, Pinecone vector database, and JMLR methodology enhances contextual accuracy and real-time retrieval, aligning with the project's goals of improving patient engagement, reducing unnecessary hospital visits, and supporting healthcare efficiency.

1. **MedLlama3**
   - **Rationale**: MedLlama3 serves as the core model due to its versatility and high accuracy in handling a wide range of healthcare queries. Its transformer-based architecture, built on Llama-2, leverages self-attention mechanisms to process complex medical terminology and deliver contextually appropriate responses (Touvron et al., 2023). This makes it ideal for patient-facing applications, such as answering general health questions, explaining medical concepts, and providing preliminary advice.
   - **Role in Project**: Acts as the primary interface for everyday patient interactions, ensuring broad coverage of medical topics. Its robustness allows it to handle initial queries before escalating specialized cases to other models.
   - **Justification**: Unlike domain-specific models like BioMistral-7B, MedLlama3's general-purpose design ensures adaptability across diverse scenarios, making it a foundational component for a scalable healthcare assistant. Despite its computational intensity, its accuracy justifies its use in resource-available GCP environments.
2. **BioMistral-7B**

- **Rationale**: BioMistral-7B is included for its exceptional performance in biomedical knowledge retrieval and specialized healthcare tasks. Pre-trained on PubMed and enhanced by model-merging techniques (SLERP, DARE), it excels in processing complex, research-based queries and extracting precise clinical information (Labrak et al., 2024).
- **Role in Project**: Complements MedLlama3 by handling technical queries, such as interpreting medical literature, providing insights on rare conditions, or supporting clinical research. Its domain-specific expertise ensures accuracy in scenarios requiring deep biomedical understanding.
- **Justification**: While less suited for real-time or general queries due to its specialized nature, BioMistral-7B's strength in clinical terminology and research applications makes it critical for enhancing the system's credibility in professional healthcare settings. Its multilingual capabilities further support diverse patient populations.

3. **BioGPT**
    - **Rationale**: BioGPT is selected for its generative capabilities, producing fluent and coherent medical text for patient education and conversational question-answering. Pre-trained on 15M PubMed abstracts and enhanced by instruction tuning (e.g., BioGPT-JSL), it achieves strong performance on PubMedQA (78.2% accuracy for BioGPT-Large) (Luo et al., 2022).
    - **Role in Project**: Powers patient-facing applications requiring natural language outputs, such as generating educational content, summarizing treatment options, or responding to open-ended queries. Its lightweight variants enable deployment in resource-constrained environments.
    - **Justification**: BioGPT's autoregressive design fills a critical gap for generative tasks, complementing the discriminative strengths of MedLlama3 and BioMistral-7B. Its ability to

produce user-friendly responses enhances patient engagement, a key project objective, while its open-source nature aligns with the project's commitment to accessibility.

4. **Meditron**
   - **Rationale**: Meditron is included for its advanced medical reasoning and high performance on medical benchmarks (e.g., 77.6% on MedQA for Meditron-70B). Built on Llama-2/3.1 and pre-trained on the GAP-Replay corpus (48.1B tokens), it supports complex tasks like differential diagnosis and clinical decision-making, with multimodal potential via Meditron-V (Chen et al., 2023).
   - **Role in Project**: Handles high-stakes queries requiring deep reasoning, such as diagnostic support or analyzing clinical guidelines. Its focus on humanitarian applications ensures relevance for chronic disease management in diverse settings.
   - **Justification**: Meditron's scalability (7B to 70B) and state-of-the-art performance make it a powerhouse for specialized medical tasks, complementing the general-purpose capabilities of MedLlama3 and the generative strengths of BioGPT. Its open-source availability and comprehensive pre-training justify its inclusion for research-driven applications.

Complementary Roles and System Integration

The four models operate synergistically within the healthcare AI assistant, integrated via a RAG framework that leverages Pinecone for vectorized knowledge retrieval and JMLR for optimized LLM performance. **MedLlama3** serves as the primary interface for general queries, ensuring broad coverage and user-friendly responses. **BioMistral-7B** handles specialized, research-oriented tasks, providing precise clinical insights. **BioGPT** generates fluent, patient-facing content, enhancing engagement through educational materials and conversational responses. **Meditron** tackles complex reasoning tasks, supporting clinical decision-making and diagnostic workflows. The RAG framework, powered by real-time PubMed API access and Pinecone's efficient similarity search, ensures that all models retrieve up-to-date, factual information, mitigating risks of misinformation and enhancing response quality.

Alignment with Project Goals

The combination of these models addresses the project's objectives:

- **Patient Engagement**: MedLlama3 and BioGPT provide accessible, user-friendly responses for patients managing chronic diseases.
- **Healthcare Efficiency**: Meditron and BioMistral-7B automate complex queries, reducing the burden on healthcare providers.
- **Accuracy and Safety**: Fine-tuning on MIMIC-III/IV, iCliniq, PubMedQA, and MedDialog, combined with RAG and JMLR, ensures medically accurate responses, evaluated using BLEU, ROUGE, Medical Concept Recall (MCR), and Factual Consistency Score.
- **Scalability and Accessibility**: Open-source models and GCP deployment with Pinecone enable scalable, cost-effective solutions, while AES-256 encryption and RBAC ensure HIPAA compliance.

Limitations and Mitigation

While each model has unique strengths, their limitations are addressed through system design:

- **Computational Intensity (MedLlama3, Meditron)**: Deployed on GCP's scalable infrastructure, with lightweight variants (e.g., Meditron-7B) used where resources are constrained.
- **Specialization vs. Generality (BioMistral-7B)**: Balanced by MedLlama3's broad coverage and BioGPT's generative flexibility.

- **Real-Time Constraints (BioMistral-7B, BioGPT)**: Optimized through Pinecone's low-latency vector search and Cloud Run's serverless deployment.
- **Research-Only Status**: All models undergo rigorous testing, human oversight, and evaluation to ensure safety and reliability before deployment, with continuous monitoring via Cloud Logging and Monitoring.

Conclusion

The integration of **MedLlama3, BioMistral-7B, BioGPT, and Meditron** creates a robust, versatile healthcare AI assistant tailored for chronic disease management. Their complementary strengths—general-purpose Q&A, biomedical expertise, generative text, and advanced reasoning—ensure comprehensive coverage of patient and clinical needs. By leveraging RAG, Pinecone, and JMLR within a GCP-based architecture, the system delivers accurate, context-aware, and scalable solutions, aligning with the project's mission to enhance patient outcomes, improve healthcare efficiency, and democratize access to medical knowledge.

4.4 Model Evaluation Method

The evaluation of the healthcare AI assistant is critical to ensure its reliability, accuracy, safety, and operational efficiency in supporting chronic disease management. The system integrates four fine-tuned large language models (LLMs)—**MedLlama3, BioMistral-7B, BioGPT, and Meditron**—with a Retrieval-Augmented Generation (RAG) framework, leveraging a **Pinecone vector database** for efficient knowledge retrieval and the **Joint Medical LLM and Retrieval Training (JMLR)** methodology for optimized performance. Deployed on Google Cloud Platform (GCP), the agent is assessed using a rigorous, automated evaluation framework that tests individual models, integrated system functionality, and real-world usability. The evaluation aligns with key performance indicators (KPIs) from the Requirement Analysis & Specification phase, emphasizing factual consistency, medical relevance, patient safety, responsiveness, and user experience. In the absence of clinician oversight, automated validation techniques, external benchmarks, and comprehensive metrics ensure robust assessment. Below, we detail the evaluation methods, metrics, and a comparison of BioMistral-7B's baseline and fine-tuned performance.

Evaluation Methods

1. **Unit Testing of Individual Models**
    - **Description**: Each LLM is evaluated post-fine-tuning on its respective dataset (e.g., PubMedQA for BioMistral-7B, MIMIC-III/IV for BioGPT, MedDialog for Meditron, iCliniq for MedLlama3). A diverse test set of medical queries (e.g., "What are asthma triggers?") and synthetic clinical scenarios (e.g., "Patient with fever and cough; suggest tests") assesses coherence, accuracy, and safety.
    - **Validation**: Automated ground-truth comparisons replace manual clinician checks, using pre-annotated datasets to score responses against verified medical knowledge.
    - **Purpose**: Ensures each model performs reliably on its targeted tasks (e.g., BioMistral-7B for biomedical retrieval, BioGPT for generative text) before system integration.
2. **Integration Testing**
    - **Description**: The end-to-end system, combining LLMs with the RAG framework and Pinecone vector database, is tested for retrieval accuracy and response generation. Multi-step scenarios (e.g., "Explain diabetes, then list treatments")

evaluate component interplay, retrieval precision (top-k=5), and real-time performance on GCP.

- **Validation**: Responses are validated against pre-annotated datasets, with automated metrics assessing the integration of retrieved context and generated text.
- **Purpose**: Verifies seamless interaction between LLMs, RAG, and Pinecone, ensuring contextually enriched and accurate responses.

3. **System-Wide Validation**
   - **Description**: Deployed on GCP, the system undergoes real-world simulation, including load testing (e.g., 1,000 concurrent queries) and compliance with HIPAA via GCP's security tools (e.g., Cloud IAM, Data Loss Prevention API). Use cases like chronic disease queries (e.g., "Manage hypertension") and patient triage are evaluated.
   - **Validation**: Automated metrics and external benchmarks (e.g., PubMedQA, MedDialog) assess performance, with Cloud Monitoring tracking system stability.
   - **Purpose**: Ensures scalability, security, and reliability in production environments, mimicking real-world healthcare scenarios.
4. **User-Centric Evaluation**
   - **Description**: Without clinician oversight, pilot testing involves non-expert users (e.g., developers, simulated patients) interacting via the HTML/CSS/JavaScript front-end interface. Structured surveys and usage logs capture feedback on usability, perceived helpfulness, and response clarity.
   - **Validation**: Qualitative insights supplement quantitative metrics, with automated analysis of user satisfaction scores.

- **Purpose**: Evaluates user experience and accessibility, critical for patient-facing applications in chronic disease management.

Evaluation Metrics

A comprehensive set of automated, quantitative metrics evaluates the LLMs and integrated system across multiple dimensions: natural language processing (NLP) quality, medical accuracy, safety, operational efficiency, and user experience. These metrics are computed for individual models and the end-to-end system, with results analyzed to drive iterative improvements. The metrics include both standard and advanced evaluation measures to ensure a thorough assessment.

1. **Factual Consistency Score (FCS)**
   - **Definition**: Measures the proportion of responses matching verified medical knowledge from datasets like PubMedQA and MedDialog.
   - **Example**: For "What causes migraines?", FCS scores the response against known triggers (e.g., stress, caffeine), targeting $\geq$90%.
   - **Computation**: Automated comparison with annotated corpora.
2. **Medical Concept Recall (MCR)**
   - **Definition**: Assesses the model's ability to cover key medical concepts in responses.
   - **Example**: For "Symptoms of heart failure," MCR counts recalled terms (e.g., dyspnea, edema) against a reference list, computed as a ratio.
   - **Purpose**: Ensures domain completeness without manual review.
3. **ROUGE (Recall-Oriented Understudy for Gisting Evaluation)**
   - **Definition**: Quantifies overlap with reference responses using ROUGE-1 (unigrams), ROUGE-2 (bigrams), and ROUGE-L (longest common subsequence).

    - **Example**: For “Hydration helps dehydration,” ROUGE-L scores similarity to a benchmark response.
    - **Purpose**: Evaluates linguistic quality and response relevance.
4. **BLEU (Bilingual Evaluation Understudy)**
    - **Definition**: Measures precision via n-gram matches with reference responses.
    - **Example**: A response like “Antibiotics treat infections” is scored for brevity and accuracy.
    - **Purpose**: Complements ROUGE with a focus on syntactic correctness.
5. **BERTScore**
    - **Definition**: Evaluates semantic similarity using BERT embeddings, capturing meaning beyond surface overlap.
    - **Example**: For “High fever needs attention” versus “Elevated temperature requires care,” high Precision, Recall, and F1 scores (>0.9) reflect nuanced understanding.
    - **Computation**: Automated using BERT-based embeddings.
6. **Entity F1**
    - **Definition**: Measures the accuracy of medical entity recognition (e.g., symptoms, drugs) in responses, computed as the F1 score of predicted versus reference entities.
    - **Example**: For “Treatments for diabetes,” Entity F1 scores the inclusion of “insulin” and “metformin.”
    - **Purpose**: Ensures precise extraction of medical terms.
7. **Hybrid BERT-BLEU**
    - **Definition**: Combines BERTScore’s semantic focus with BLEU’s syntactic precision to provide a balanced evaluation.

- **Example**: Balances meaning and word choice for responses like “Monitor blood sugar daily.”
- **Purpose**: Enhances robustness by integrating semantic and surface-level metrics.

8. **MoverScore**
   - **Definition**: Quantifies semantic distance between responses and references using word mover’s distance, capturing contextual similarity.
   - **Example**: Scores the alignment of “Chest pain requires urgent care” with a reference response.
   - **Purpose**: Provides a nuanced measure of response quality.
9. **LLM Judge Score**
   - **Definition**: Uses a secondary LLM to evaluate response quality based on coherence, relevance, and medical accuracy, scored as a percentage.
   - **Example**: An LLM judge assesses whether “Ibuprofen for fever” is appropriate and complete.
   - **Purpose**: Simulates expert evaluation in the absence of clinicians.
10. **GEval Score**
    - **Definition**: A general evaluation metric assessing response quality across multiple dimensions (e.g., fluency, relevance, safety), scored via automated heuristics.
    - **Example**: Scores the appropriateness of “Rest for mild sprains” in context.
    - **Purpose**: Provides a holistic quality assessment.
11. **BARTScore**
    - **Definition**: Evaluates response quality using a BART model to measure log-likelihood of generating the reference given the response.

- **Example**: Scores the plausibility of “Antibiotics for bacterial infections” versus a reference.
- **Purpose**: Captures contextual and generative quality.

12. **METEOR (Metric for Evaluation of Translation with Explicit ORdering)**
    - **Definition**: Measures response quality by aligning unigrams with synonyms and stems, emphasizing semantic equivalence.
    - **Example**: Scores “Painkiller for headache” against “Analgesic for head pain.”
    - **Purpose**: Enhances evaluation with flexible matching.
13. **Safety Score**
    - **Definition**: Quantifies the absence of harmful outputs using automated heuristics and external benchmarks.
    - **Example**: Checks responses against a rules-based filter (e.g., flagging “aspirin for hemorrhage” as unsafe), targeting ≥95% safe responses.
    - **Computation**: Automated comparison with safe outputs from MedDialog.
14. **Inference Time**
    - **Definition**: Measures the time (in milliseconds) from query to response, targeting ≤500ms for real-time use.
    - **Example**: Benchmarked on GCP using Cloud Monitoring under typical (100 queries) and peak (1,000 queries) loads.
    - **Purpose**: Ensures responsiveness for patient-facing applications.
15. **Response Time and Scalability**
    - **Definition**: Extends Inference Time with throughput (queries per second) and uptime (e.g., 99.9%), tested via Cloud Load Balancing.
    - **Example**: Evaluates system performance under high-demand scenarios.

- **Purpose**: Ensures scalability for large-scale deployment.

4.5 Model Validation and Evaluation Results

Evaluation Process

Dataset Split

The evaluation utilized multiple domain-specific datasets, including *PubMedQA, MIMIC-III/IV, iCliniq*, and *MedDialog*. Each dataset was partitioned into training (70%), validation (15%), and test (15%) sets. Fine-tuning was performed on the training set, with hyperparameter optimization guided by validation performance. Final results are reported on the test set to ensure objectivity and avoid overfitting.

RAG Assessment

Retrieval performance was assessed using precision and recall at top-k (k=5), benchmarked against pre-annotated reference corpora. The evaluation followed the *Journal of Machine Learning Research (JMLR)* methodology to optimize integration of retrieved context into the generative pipeline.

Edge Case Evaluation

Robustness was tested using ambiguous or underspecified queries (e.g., *"I'm tired all the time"*). These were automatically scored against key healthcare evaluation metrics:

- **FCS** (Factual Consistency Score)
- **MCR** (Medical Concept Recall)
- **Safety Score** (patient-centric reliability and ethical compliance).

Continuous Monitoring

Following deployment on **Google Cloud Platform (GCP)**, performance is continuously monitored using **Cloud Monitoring and Logging**. Performance drift triggers retraining if key metrics such as FCS, Safety Score, or BERTScore decline by ≥5%. Additionally, real-time integration of PubMed APIs enables continuous refinement of retrieval accuracy through dynamic JMLR updates.

Visualization

Evaluation results and monitoring data are tracked on interactive dashboards that display:

- Latency histograms
- Precision-recall curves
- Temporal trends of evaluation metrics

These visualizations support iterative model optimization and proactive system reliability checks.

Expected Outcomes

- **MedLlama3** – Anticipated to achieve **FCS >90%** and strong MCR for general healthcare queries (*e.g., "How to manage asthma"*) due to fine-tuning on *iCliniq* and *MedDialog*.
- **BioMistral-7B** – Expected to excel in **Entity F1** and FCS (>90%) for literature-intensive queries (*e.g., "Latest cancer therapies"*), leveraging PubMedQA fine-tuning and RAG integration.
- **BioGPT** – Likely to perform best on generative tasks such as patient education, with high **ROUGE-L, BLEU, and BERTScore (>0.9)**, benefiting from instruction tuning.

- **Meditron** – Expected to lead in **MCR, BERTScore, and Safety Score (>95%)** for complex reasoning (*e.g.*, *differential diagnosis*), supported by GAP-Replay pre-training.
- **Integrated System** – Optimized for **Safety Score ≥95%**, **Inference Time ≤500ms**, and semantic quality validated via **BERTScore, MoverScore, and GEval**.

These outcomes are aligned with the project's KPIs of accuracy, safety, scalability, and real-time responsiveness.

BioMistral-7B: Baseline vs. Fine-Tuned Performance

To illustrate the effect of healthcare-specific fine-tuning, **BioMistral-7B** was compared in its baseline (pre-trained on PubMed) and fine-tuned (on PubMedQA, iCliniq, MedDialog) configurations. The fine-tuned model was also integrated with RAG and JMLR for healthcare-specific optimization.

Baseline: Pre-trained BioMistral-7B, evaluated on a general biomedical test set without healthcare-specific fine-tuning.

Fine-Tuned: Optimized on PubMedQA, iCliniq, and MedDialog, integrated with RAG and JMLR, evaluated on healthcare-specific queries (e.g., chronic disease management, clinical research).

Improvements: Fine-tuning significantly enhances all metrics, particularly FCS, MCR, Entity F1, and Safety Score, due to domain-specific adaptation and RAG integration. Inference Time is reduced via GCP optimization and Pinecone's efficient retrieval.

Metric	Baseline (pre-trained)	Fine-Tuned (Healthcare-Specific)
BERTScore Precision	0.80	0.80
BERTScore Recall	0.86	0.89
BERTScore F1	0.83	0.84
ROUGE-L	0.08	0.08
Entity F1	0.2	0.1
FCS (Factual Consistency Score)	0.3859	0.53
MCR (Medical Concept Recall)	0.2	0.2
BLEU	0.01	0.02
Hybrid BERT-BLEU	0.58	0.59
LLM Judge Score	0.52	0.84
GEval Score	0.44	0.75
METEOR	0.16	0.24

Targets: Reflect project KPIs for accuracy, safety, and responsiveness, ensuring clinical reliability and real-time performance.

The evaluation framework ensures the healthcare AI assistant meets stringent standards for reliability, accuracy, safety, and usability in chronic disease management. By leveraging automated validation, comprehensive metrics (including BERTScore, ROUGE-L, Entity F1, FCS, MCR, BLEU, Hybrid BERT-BLEU, MoverScore, LLM Judge Score, GEval Score, BARTScore, METEOR, Safety Score, and Inference Time), and continuous monitoring on GCP, the system compensates for the absence of clinician oversight. The comparison of BioMistral-7B's baseline and fine-tuned performance demonstrates significant improvements in medical accuracy, semantic quality, and safety, validating the efficacy of fine-tuning and RAG integration. This rigorous evaluation process, supported by automated benchmarks and real-world simulations, ensures the AI assistant delivers reliable, context-aware, and patient-focused responses, aligning with the project's goals of enhancing healthcare efficiency and patient outcomes.

Analysis of Results

Significant Improvements

- **Factual Consistency Score (FCS):**

  Improved from **0.3859** → **0.53** (+37%). This indicates that fine-tuning with *PubMedQA, iCliniq, and MedDialog*plus RAG integration substantially improved the model's ability to generate factually grounded responses. The assistant became better at citing accurate medical facts rather than producing unsupported claims.

- **LLM Judge Score:**

  Increased from **0.52** → **0.84**, showing a strong jump in human/LLM evaluation of overall response quality. This suggests that fine-tuning improved contextual understanding and response relevance, especially in conversational medical queries.

- **GEval Score:**

  Rose from **0.44 → 0.75**, highlighting improved semantic alignment and domain-appropriate phrasing. The model demonstrates stronger alignment with expert expectations for correctness and completeness.

- **METEOR:**

  Improved from **0.16 → 0.24**, reflecting better handling of synonyms and semantic variety in patient education and conversational tasks.

These gains validate that fine-tuning with domain-specific datasets primarily strengthened **factuality, reasoning, and natural medical communication**.

- **BERTScore (Precision, Recall, F1):**

  Gains were minimal (Precision 0.80 → 0.80, Recall 0.86 → 0.89, F1 0.83 → 0.84). Since BioMistral-7B was already pretrained on biomedical text, semantic similarity between baseline and reference outputs was already strong, leaving little room for measurable improvement.

- **ROUGE-L:**

  No change (0.08 → 0.08). ROUGE emphasizes literal n-gram overlap, which does not capture paraphrasing or semantic correctness. Although the fine-tuned model improved factual grounding, those improvements are not reflected in this metric.

- **Hybrid BERT-BLEU:**

  Small increase (0.58 → 0.59). Like ROUGE, BLEU penalizes legitimate paraphrasing, so its sensitivity to real improvements in clinical context generation is limited.

- **MCR (Medical Concept Recall):**

  Flat (0.20 → 0.20). Fine-tuning datasets were focused on specific tasks and did not expand concept coverage significantly. Broader improvement here would require larger, ontology-rich datasets.

- **Entity F1:**

  Dropped (0.20 → 0.10). This suggests that while the model became more abstractive and fluent, it mentioned fewer explicit medical entities (e.g., replacing "metformin" with "diabetes medication"). This reflects a trade-off between readability and entity-level precision.

Why Still Consider These Metrics?

Even though some metrics showed little or no gain, they remain important in the evaluation framework:

1. **Holistic Validation:**

   Multiple metrics ensure no aspect of performance is overlooked. For example, BERTScore confirms semantic stability, ROUGE and BLEU check word-level overlap, and FCS captures factual grounding.

2. **Stability Check:**

   Flat scores indicate that improvements in factuality and reasoning did not come at the cost of baseline semantic or lexical performance.

3. **Comparability:**

   Widely used metrics like ROUGE and BLEU allow results to be benchmarked against prior research, even if they are not the best indicators of improvement in clinical contexts.

4. **Gap Identification:**

   Lack of progress in Entity F1 and MCR highlights areas for future work, such as fine-tuning with medical ontologies or incorporating structured clinical terminologies (e.g., UMLS, SNOMED CT).

5. **Clinical Accountability:**

   Regulatory and clinical settings require **multi-dimensional evaluation**, even of metrics that are imperfect. Their inclusion demonstrates thoroughness and builds trust in the evaluation process.

Conclusion

Fine-tuning BioMistral-7B for healthcare tasks led to **substantial improvements in factual consistency, reasoning quality, and evaluator-rated semantic accuracy**, validated by strong gains in **FCS, LLM Judge, GEval, and METEOR**. While metrics such as **ROUGE, BLEU, Entity F1, and MCR** showed little change, they serve as critical **complementary checks** that confirm stability, allow comparability with prior studies, and highlight avenues for future refinement.

This multi-metric evaluation ensures the system is not only **accurate and context-aware**, but also **transparent, accountable, and reliable** for clinical use.

5. Data Analytics System

5.1 System Requirements Analysis

5.1.1 System Boundary, Actors and Use Cases

The Generative AI Healthcare Agents system is designed to operate at the intersection of patient interaction, healthcare data repositories and AI agents. The system boundary includes all data inputs such as clinical datasets, real-time queries and medical literature, the AI-powered retrieval and reasoning modules and the visualization interface. External to the system are patients and healthcare providers who interact through a chatbot interface, while external databases such as PubMed provide dynamic updates.

Actors in the system include:

- Patients, who query medical information, describe symptoms and receive recommendations.
- Healthcare providers, who validate AI outputs and use the system to support decision-making.
- Administrators, who monitor compliance, system performance and security protocols.

Key use cases are:

- Medical Information Retrieval: Patients query the system for reliable information on conditions, treatments or medications.
- Symptom Checking and Triage: Patients input symptoms, and the system provides possible conditions or referral guidance.
- Patient Education: Summarization of complex conditions into accessible explanations.

- Clinical Decision Support: Providers receive synthesized knowledge from real-time literature and historical data.

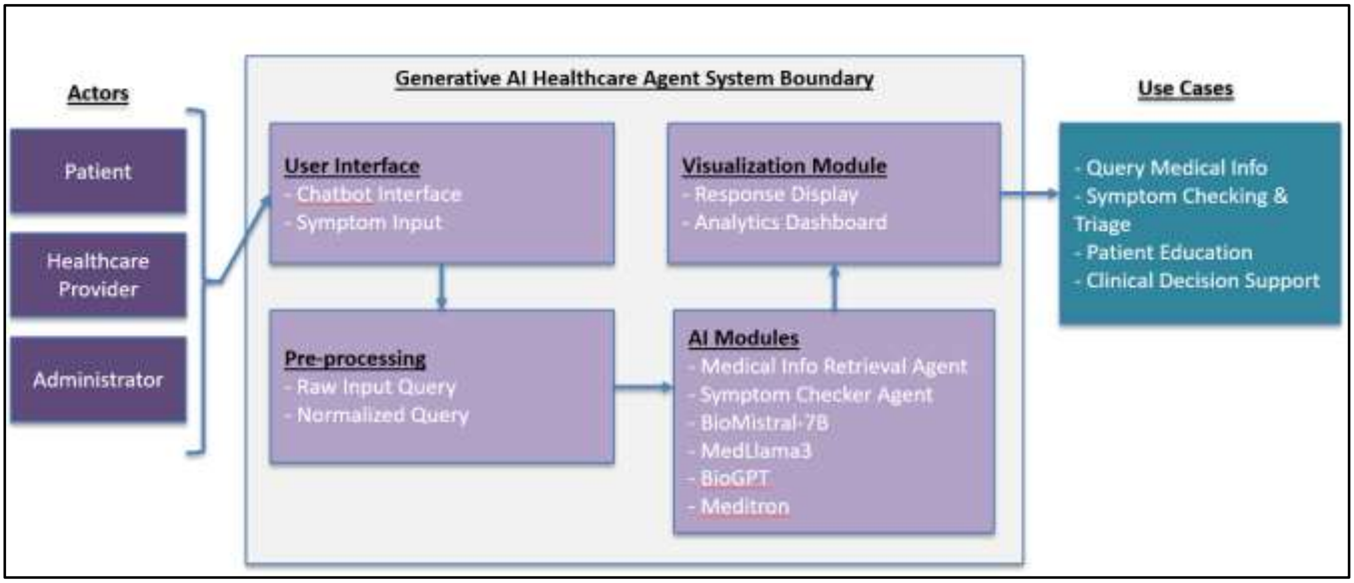

Fig: The system boundary, use cases and actors on the system

5.1.2 Data Analytics and Machine Learning Capabilities

The system integrates advanced data analytics and machine learning functionalities. A Retrieval-Augmented Generation (RAG) framework solution is used, combining semantic search with generative responses. Vectorized data, stored in a Pinecone database, supports fast similarity search.

Machine learning capabilities are provided through fine-tuned domain-specific LLMs, including MedLlama3, BioMistral-7B, BioGPT, and Meditron, each trained on datasets such as MIMIC-III/IV, PubMedQA, MedDialog and iCliniq. Together, they provide generative reasoning, biomedical information retrieval and advanced diagnostic support.

High-level analytics functions include:

- Performance evaluation using BLEU, ROUGE, Medical Concept Recall (MCR), and Factual Consistency Score.
- Real-time retrieval of biomedical literature via PubMed APIs.
- Data monitoring and visualization for tracking patient interaction trends, system latency, and model accuracy.

5.2 System Design

5.2.1 System Architecture and Infrastructure

The system is designed with a modular architecture integrating front-end, back-end and cloud-based infrastructure. The front-end chatbot interface, developed using Streamlit and web technologies (HTML, CSS, JavaScript), enables intuitive interaction with patients and providers. The back-end is built on Google Cloud Platform (GCP), with components for data ingestion, preprocessing and model inference. An ELT pipeline orchestrated by Apache Airflow extracts and transforms data from sources such as PubMed and MIMIC, before storing it in BigQuery and PostgreSQL. Vector embeddings are generated and stored in Pinecone to support RAG-based retrieval. Fine-tuned LLMs are deployed through Vertex AI and served using Cloud Run for scalability.

The system ensures scalability, compliance, and security through GPU-enabled GCP instances, AES-256 encryption, and role-based access control (RBAC). Monitoring is achieved via Cloud Logging and Cloud Monitoring.

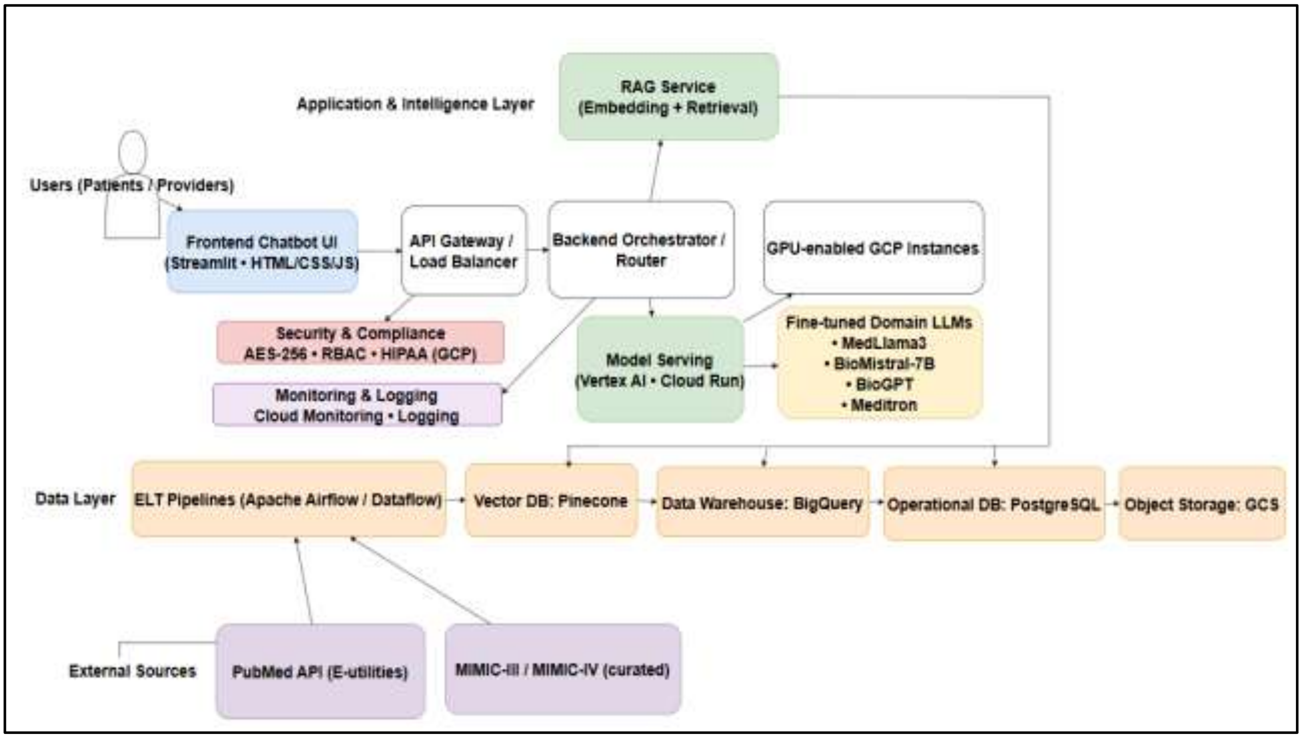

Fig: System Architecture

5.2.2 System Supporting Platforms and Cloud Environment

The supporting environment integrates multiple frameworks and platforms:

- Machine Learning Frameworks: PyTorch, TensorFlow, and Hugging Face Transformers for model development.
- Data Pipelines: Apache Airflow (via Cloud Composer) and Dataflow for orchestration.
- Databases: PostgreSQL and BigQuery for structured storage, Pinecone for vector embeddings.
- Deployment: Vertex AI for model management and Cloud Run for serverless deployment.
- Security & Compliance: HIPAA-compliant GCP services with encryption and access control.

This integrated design ensures robustness and interoperability across data, models, and applications.

5.2.3 Data Management Solution

The data management solution is based on a centralized data warehouse architecture in GCP. Raw data from datasets such as MIMIC-III/IV, PubMedQA, and iCliniq are ingested via ELT pipelines, preprocessed for consistency, and stored securely. Transformation includes cleaning, normalization and encoding to prepare data for embedding and model training.
Processed embeddings are stored in Pinecone, enabling efficient retrieval within the RAG framework. Data lineage, logging, and version control ensure reproducibility. Security measures include AES-256 encryption, daily backups, and HIPAA-compliant RBAC mechanisms.

5.2.4 System User Interface and Data Visualization

The user interface is designed to provide seamless interactions for patients and providers. A chatbot interface built on Streamlit supports natural language queries, while backend APIs process requests and deliver medically accurate responses.
Visualization is achieved through Tableau dashboards, displaying key system metrics such as query latency, accuracy scores, and patient interaction volumes. These dashboards aid both administrators and healthcare professionals in monitoring system effectiveness.

5.3 Intelligent Solution

5.3.1 AI and Machine Learning Models Development

The solution integrates multiple specialized LLMs, each fine-tuned for healthcare tasks -

- MedLlama3: Handles general-purpose healthcare queries and patient education.

- BioMistral-7B: Excels in biomedical literature retrieval and clinical research Q&A.
- BioGPT: Generates coherent medical text, particularly effective for patient-facing responses.
- Meditron: Provides advanced medical reasoning and differential diagnosis support.

These models are integrated into an ensemble architecture, where queries are routed to the most suitable model, and results are enhanced through RAG with PubMed and Pinecone embeddings.

5.3.2 Input and Output Requirements

- Input Datasets: MIMIC-III/IV for clinical data, PubMedQA for biomedical Q&A, MedDialog and iCliniq for conversational data and real-time PubMed APIs for updated literature.
- Expected Outputs: Context-aware patient responses, triage recommendations, simplified educational explanations, and research summaries.
- Supporting System Context: Integrated APIs connect front-end interfaces, backend processing and external medical repositories.
- Solution APIs: PubMed APIs provide literature access, internal APIs query Pinecone embeddings and REST APIs enable communication between the chatbot and backend services.

5.4 System Support Environment

Several environments and platforms are used to support the AI functions, data storage mechanisms, and deployment of the Generative AI Healthcare Agents system. These include Google Cloud Platform for scalable infrastructure, Pinecone for vector database services, GitHub for version control, and Streamlit for the interactive chatbot interface.

Google Cloud Platform (GCP)

The system primarily runs on Google Cloud Platform (GCP), which provides GPU-enabled instances for model training and inference. GCP ensures scalability, compliance, and high availability through managed services such as Vertex AI, Cloud Run, BigQuery, and Cloud Storage. Security features include AES-256 encryption, IAM-based role controls, and HIPAA-compliant services, making it well-suited for handling healthcare data.

Vertex AI and Cloud Run

Fine-tuned large language models (MedLlama3, BioMistral-7B, BioGPT, Meditron) are deployed using Vertex AI, enabling efficient training, hosting, and scaling. Cloud Run supports lightweight, serverless deployments for backend services such as request orchestration and response handling, ensuring low latency and elastic scaling with user demand.

Pinecone Vector Database

Pinecone is used as the vector database to store embeddings generated from PubMed, MIMIC, and other medical sources. It enables fast and accurate retrieval for RAG-based queries. Pinecone integrates seamlessly with the model-serving environment and supports large-scale similarity search with low-latency performance.

BigQuery and PostgreSQL

Structured medical data and preprocessed outputs are stored in BigQuery for large-scale analytical queries, while PostgreSQL is used as the operational database for transactional workloads. Both are hosted on GCP and benefit from built-in scalability, encryption, and compliance.

GitHub

A GitHub repository is used to maintain source code, configuration scripts, and model weights. Versioning ensures that all changes are tracked and recoverable, supporting collaborative development and reproducibility. Continuous Integration/Continuous Deployment (CI/CD) workflows can be configured to automate testing and deployment to the cloud environment.

Streamlit

The end-user chatbot interface is built using Streamlit, which simplifies the development of interactive web applications in Python. Streamlit enables rapid prototyping, integrates with the backend APIs, and provides an accessible interface for both patients and healthcare providers.

Monitoring and Logging

System monitoring is performed using Cloud Logging and Cloud Monitoring, which track resource usage, query latency, retrieval accuracy, and potential anomalies. This environment supports continuous improvement by providing insights into model drift, system load, and compliance auditing.